

Introduction to IPCluster

Sponsored by AWS

Continuum Analytics



Unofficial Python Slogan

*“CPU time is cheap, **my** time is expensive”*



Some Truthy Statements

- We prefer an expressive language over raw speed



Some Truthy Statements

- We prefer an expressive language over raw speed
 - 97% of the time, it's fast enough every time
 - Premature optimization is the root of all evil



Some Truthy Statements

- We prefer an expressive language over raw speed
 - 97% of the time, it's fast enough every time
 - Premature optimization is the root of all evil
- But scientific workflows can get big



Some Truthy Statements

- We prefer an expressive language over raw speed
 - 97% of the time, it's fast enough every time
 - Premature optimization is the root of all evil
- But scientific workflows can get big
 - My expensive time spent waiting instead of doing science
 - Let's replace that with more cheap CPU time



Some Truthy Statements

- We have horsepower on-demand (AWS, Azure, Heroku, University Clusters)




Some Truthy Statements

- We need a nice way to wrangle the data/code between nodes



IPCluster Overview



The Basic Use Case

- An embarrassingly (pleasingly) parallel workflow
 - Little/no communication between jobs
 - Parameter scans, multiple BLAST queries, genetic algorithms, etc.



The Basic Use Case

- An embarrassingly (pleasingly) parallel workflow
 - Little/no communication between jobs
 - Parameter scans, multiple BLAST queries, genetic algorithms, etc.
- Can utilize the CPU/GPU/RAM of many nodes simultaneously



The Basic Use Case

- An embarrassingly (pleasingly) parallel workflow
 - Little/no communication between jobs
 - Parameter scans, multiple BLAST queries, genetic algorithms, etc.
- Can utilize the CPU/GPU/RAM of many nodes simultaneously
- Moves relatively small amounts of data
 - Aggregates/subsets are good enough
 - The data is already on the nodes



Parallel-o World

Create a client object

```
from IPython.parallel import Client  
  
client = Client()
```



Parallel-o World

Create a client object

```
from IPython.parallel import Client  
  
client = Client()
```

Construct a view

```
balanced = client.load_balanced_view()  
# or  
direct = client[:]
```



Parallel-o World

Define a function

```
def hello():  
    return "Hello World!"
```



Parallel-o World

Define a function

```
def hello():  
    return "Hello World!"
```


Run it remotely

```
resp = direct.apply(hello)  
resp.get()
```

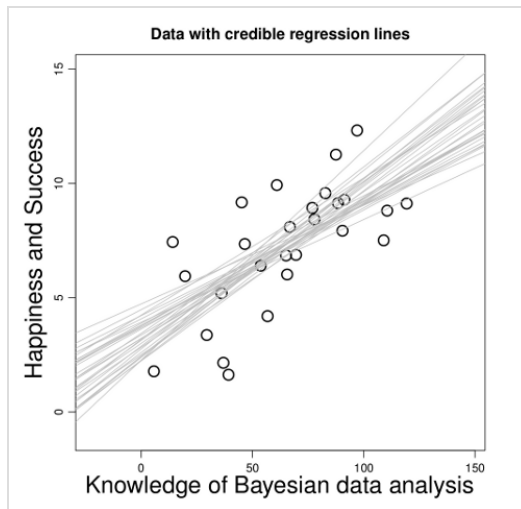
```
Out[33]: ['Hello World!', 'Hello World!', 'Hello World!']
```



Easy Demo




Use Case: Bayesian Estimation




Images: John K Kruschke

Use Case: Bayesian Estimation



Use Case: Bayesian Estimation



Bayesian Demo

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$